

Chapter 26: Concurrency - An Introduction

Overview

- **Main Idea:** Threads enable multiple points of execution within a single process, sharing the same address space but having private registers/stacks.
- **Key Terms:** Thread, Process, Address Space, Context Switch, Race Condition, Critical Section, Mutual Exclusion.

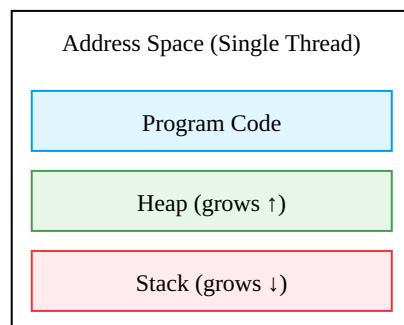
1. Threads vs Processes

Key Differences

| Feature | Thread | Process |
|----------------|-------------------------------|---------------------------|
| Address Space | Shared with other threads | Private |
| Context Switch | Faster (no page table switch) | Slower |
| Communication | Direct (shared memory) | Complex (pipes, messages) |

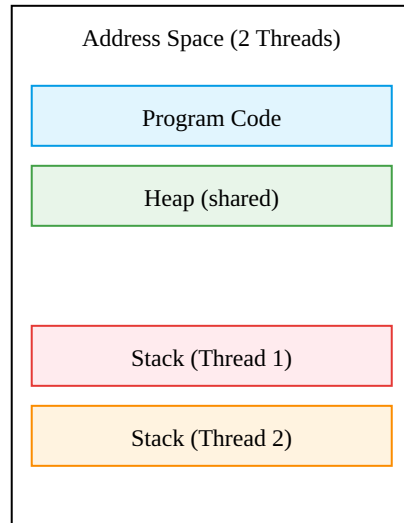
Address Space Layout

Single-Threaded Process



(Stack grows downward, heap upward.)

Multi-Threaded Process (2 threads)



*(Each thread has its own stack; heaps/code are shared.)__

2. Why Use Threads?

Parallelism Example

```
// Parallel array addition (pseudo-code)
#pragma omp parallel for
for (int i = 0; i < N; i++) {
    C[i] = A[i] + B[i]; // Threads split work
}
```

I/O Overlap

```
// Web server handling concurrent requests
while (1) {
    int conn_fd = accept(); // Blocking call
    pthread_create(&thread, NULL, handle_request, (void*)conn_fd);
}
```

3. Thread Creation Code

```
#include <pthread.h>
void *mythread(void *arg) {
```

```

    printf("%s\n", (char *)arg); // Thread-local work
    return NULL;
}

int main() {
    pthread_t t1, t2;
    pthread_create(&t1, NULL, mythread, "A"); // Create Thread 1
    pthread_create(&t2, NULL, mythread, "B"); // Create Thread 2
    pthread_join(t1, NULL); // Wait for threads
    pthread_join(t2, NULL);
    printf("main: end\n");
}

```

Possible Execution Traces

1. Order 1:

```

main      - - - - _creates T1, T2 - - - - -
Thread 1 (A) _____ prints "A" _____
Thread 2 (B) _____ prints "B" _____

```

(main → A → B)

2. Order 2:

```

main      - - - - _creates T1, T2 - - - - -
Thread 1 (A) _____ prints "A" _____
Thread 2 (B) _____ prints "B" _____

```

(main → B → A)

4. Race Conditions

Problematic Code

```
static volatile int counter = 0;
void *mythread(void *arg) {
    for (int i = 0; i < 1e7; i++) {
        counter++; // Race condition!
    }
}
```

- **Expected:** 20,000,000 (2 threads × 10M increments).
- **Actual:** Random (e.g., 12,345,678) due to interleaving.

Assembly Breakdown

```
mov 0x8049a1c, %eax ; Load counter
add $0x1, %eax      ; Increment (race here!)
mov %eax, 0x8049a1c ; Store back
```

| | | | |
|----------|-------------------|---------------|-------------------|
| Thread 1 | mov counter, %eax | add \$1, %eax | mov %eax, counter |
| Thread 2 | mov counter, %eax | add \$1, %eax | mov %eax, counter |

Final counter = 51 (expected 52)

*(Interrupts corrupt final value.)__

5. Critical Sections & Solutions

Atomic Operation Wish

```
memory-add 0x8049a1c, $0x1 ; Hypothetical atomic increment
```

Synchronization Primitives

- **Mutexes:** pthread_mutex_lock(&lock)
 - **Semaphores:** sem_wait(&sem)
-

6. Key Concepts

| Term | Definition |
|-------------------------|---|
| Critical Section | Code accessing shared data (e.g., <code>counter++</code>). |
| Race Condition | Undefined behavior due to unsynchronized access. |
| Mutual Exclusion | Ensuring only one thread executes critical section at a time. |

7. Why OS Cares

- **Kernel Data Structures:** Threads manage concurrent access to OS resources (e.g., file systems).
- **Synchronization:** Built via hardware (atomic ops) + OS (scheduler).

Homework Notes

1. `loop.s`:

```
./x86.py -t 1 -p loop.s -i 100 -R dx # Single-threaded
./x86.py -t 2 -p loop.s -i 3 -r      # Race with randomness
```

2. Critical Section Identification:

Critical Section
`mov counter, %eax`
`add $1, %eax`

`mov %eax, counter`
`*(Interrupts must avoid counter++ region.)__`

References

- Dijkstra, E. W. (1965). *Solution of a problem in concurrent programming control*.
- Gray, J. (1992). *Transaction Processing: Concepts and Techniques*.