

Chapter 27: Interlude - Thread API

Overview

- **Main Idea:** POSIX threads (pthreads) provide interfaces for thread creation, synchronization (locks/condition variables), and management.
 - **Key Terms:** pthread_create, pthread_join, mutex, condition variable, critical section.
-

1. Thread Creation


pthread_create Signature

```
#include <pthread.h>
int pthread_create(
    pthread_t *thread,           // Thread identifier
    const pthread_attr_t *attr,  // Thread attributes (NULL for defaults)
    void *(*start_routine)(void *), // Function pointer to thread's entry point
    void *arg                    // Argument passed to start_routine
);
```

Example: Passing Arguments

```
typedef struct { int a; int b; } myarg_t;
void *mythread(void *arg) {
    myarg_t *args = (myarg_t *) arg;
    printf("%d %d\n", args->a, args->b);
    return NULL;
}

int main() {
    pthread_t p;
    myarg_t args = { 10, 20 };
    pthread_create(&p, NULL, mythread, &args); // Pass struct as argument
    pthread_join(p, NULL);
}
```



(Main thread creates worker thread with arguments.)

2. Thread Completion

pthread_join Signature

```
int pthread_join(pthread_t thread, void **retval); // Waits for thread to finish
```

Example: Returning Values

```
typedef struct { int x; int y; } myret_t;
void *mythread(void *arg) {
    myret_t *rvals = malloc(sizeof(myret_t));
    rvals->x = 1;
    rvals->y = 2;
    return (void *) rvals; // Return heap-allocated struct
}

int main() {
    pthread_t p;
    myret_t *rvals;
    pthread_create(&p, NULL, mythread, NULL);
    pthread_join(p, (void **) &rvals);
    printf("Returned %d %d\n", rvals->x, rvals->y);
    free(rvals);
}
```

Warning: Never return stack-allocated data from a thread!

```
// BAD CODE: Returning stack address
myret_t oops;
return (void *) &oops; // Crash-prone!
```


3. Locks (Mutexes)

Initialization

```
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER; // Static
// OR
pthread_mutex_init(&lock, NULL); // Dynamic
```

Usage

```
pthread_mutex_lock(&lock);
x = x + 1; // Critical section
pthread_mutex_unlock(&lock);
```



*(Threads serialize access to critical sections.)__

Common Pitfalls

1. Missing Initialization:

```
pthread_mutex_t lock; // Uninitialized + undefined behavior!
pthread_mutex_lock(&lock);
```

2. Ignoring Error Codes:

Always check return values or use wrappers:

```
void Pthread_mutex_lock(pthread_mutex_t *mutex) {
    int rc = pthread_mutex_lock(mutex);
    assert(rc == 0);
}
```

4. Condition Variables


Initialization

```
pthread_cond_t cond = PTHREAD_COND_INITIALIZER;
```

Waiting and Signaling

```
// Thread 1: Waits for condition
pthread_mutex_lock(&lock);
while (ready == 0) {
    pthread_cond_wait(&cond, &lock); // Releases lock while waiting
}
pthread_mutex_unlock(&lock);
```

```
// Thread 2: Signals condition
pthread_mutex_lock(&lock);
ready = 1;
pthread_cond_signal(&cond);
pthread_mutex_unlock(&lock);
```



(Signaling thread wakes up waiting thread.)

Key Points:

- Always use a **while** loop (not **if**) for conditions due to spurious wakeups.
- Hold the mutex when modifying shared state (e.g., **ready**).

5. Thread API Guidelines

Do	Don't
Initialize locks/condition variables.	Use uninitialized synchronization primitives.
Check return codes.	Ignore error handling.
Pass heap-allocated data to threads.	Return stack addresses from threads.
Use condition variables for signaling.	Spin-wait on flags (while (flag == 0);).

Homework Insights

1. Race Conditions:

```
valgrind --tool=helgrind main-race # Detects data races
```

2. Deadlocks:

Circular lock acquisition (e.g., T1 holds A, waits for B; T2 holds B, waits for A).

3. Efficiency:

Spin-waiting (**while (done == 0);**) wastes CPU; condition variables are preferred.