

Chapter 13: The Abstraction: The Address Space

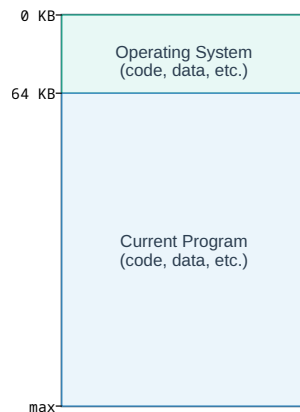
This chapter introduces the **address space**, the fundamental abstraction the Operating System provides for managing memory. It explains why this abstraction is necessary and what goals it aims to achieve, setting the stage for understanding how memory is virtualized.

1. The Evolution of Memory Management

1.1. Early Systems: Direct Physical Memory Access

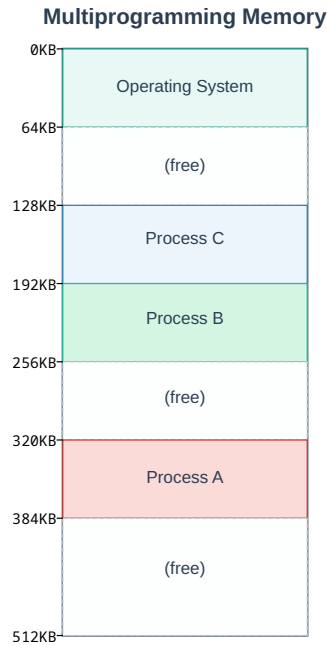
- In the earliest computer systems, there was no abstraction. The OS was a simple library residing at a fixed physical memory location (e.g., address 0), and a single running program used the rest of the physical memory directly.
- This was simple but offered no protection and could only run one program at a time.

Early System Memory Layout



1.2. The Rise of Multiprogramming and Time Sharing

- To improve the utilization of expensive machines, **multiprogramming** was introduced. The OS would load multiple processes into physical memory simultaneously.
- When one process performed a slow I/O operation, the OS could switch the CPU to another process, keeping the CPU busy.
- This led to **time sharing**, where users could interact with the system concurrently, demanding timely responses.
- To make time sharing efficient, processes had to reside in memory at the same time. Swapping entire processes to and from disk for every context switch was far too slow.



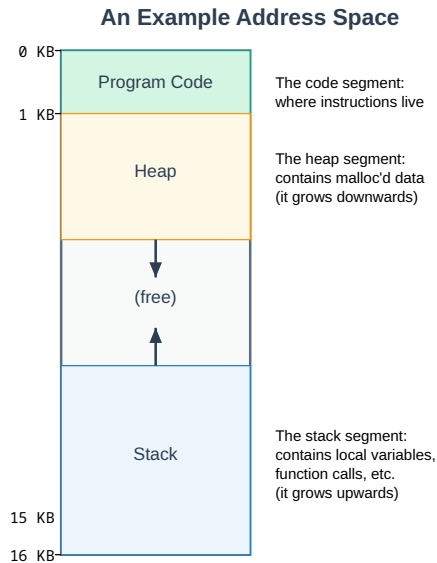
- **New Challenge:** With multiple processes in memory, **protection** becomes critical. The OS must prevent one process from reading or writing the memory of another process.
-

2. The Abstraction: The Address Space

To solve the problems of protection, ease of use, and efficiency, the OS provides each process with its own **address space**.

- **Definition:** The address space is the running program's *view* of memory. It is the abstraction of physical memory created by the OS for each process.
- **Illusion of Privacy:** Each process believes it has its own large, private memory to work with. It is completely isolated from the memory of other processes.
- **Structure of an Address Space:** The address space contains all of a program's memory state. It is typically organized into three main segments:
 1. **Code:** The program's executable instructions. This segment is static and usually read-only.
 2. **Stack:** Used for local variables, function parameters, and return addresses. It grows and shrinks as functions are called and return. By convention, it starts at the high end of the address space and grows downwards (towards lower addresses).
 3. **Heap:** Used for dynamically allocated memory (e.g., via `malloc()`)

or **new**). It is managed by the programmer and, by convention, grows upwards (towards higher addresses).



3. The Core Problem: Virtualizing Memory


The address space is an *abstraction*. The program sees a neat, contiguous block of memory starting at address 0, but in reality, the OS has placed its data at arbitrary locations in physical memory.

The Crux of the Problem: How to Virtualize Memory?

How can the OS build this abstraction of a private, potentially large address space for multiple running processes (all sharing memory) on top of a single, physical memory?

- **Virtual vs. Physical Addresses:**
 - An address generated by a running program is a **virtual address**. It exists only within the context of that process's address space.
 - The OS, with hardware assistance, translates this virtual address into a **physical address** before accessing main memory.
- **Example:** A process might issue a **load** instruction for virtual address 0x1000. The OS and hardware translate this to the correct physical address where that data is actually stored (e.g., 0x50000).

Virtual-to-Physical Address Translation



4. Goals of a Virtual Memory System

A modern Virtual Memory (VM) system is designed with three primary goals in mind:

1. **Transparency:** The memory virtualization should be invisible to the application. The program should behave as if it has its own private, dedicated physical memory, without being aware of the underlying translation mechanisms.
2. **Efficiency:** The virtualization process must be efficient in both time and space.
 - **Time:** Address translations must be fast to avoid slowing down the program. This requires significant hardware support (e.g., TLBs).
 - **Space:** The data structures used by the OS to manage virtualization should not consume excessive memory.
3. **Protection:** The OS must protect processes from one another and protect itself from user processes.
 - **Isolation:** This is the key principle. By ensuring each process can only access memory within its own address space, the OS isolates processes, preventing one from interfering with another or with the kernel.

5. Summary

The **address space** is the OS's fundamental abstraction of memory, providing each process with the illusion of a large, private, and contiguous memory space. This illusion is built upon the mechanism of **virtual memory**, where the OS and hardware work together to translate **virtual addresses** generated by the program into actual **physical addresses**. This system is designed to be transparent, efficient, and secure, providing crucial protection and isolation between processes.